Unsupervised Machine Learning for Autonomous Discovery of Thermodynamic Phase Transitions in the 2D Ising Model

Harshiddh Parmar

Indian Institute of Science Education and Research (IISER) Pune

Abstract: Can an artificial intelligence discover physical laws just by looking at raw data? In this project, we simulate the 2D Ferromagnetic Ising Model to see if unsupervised machine learning can figure out a phase transition on its own. Instead of feeding the AI traditional physics equations, we gave it thousands of raw "snapshots" of the spin grid. Using Principal Component Analysis (PCA) and K-Means clustering, the AI successfully grouped the data into distinct physical phases without any temperature labels. Even better, the algorithm's first principal component perfectly mirrored the system's actual macroscopic magnetization. By analyzing how the data fluctuated over time, we also proved the divergence of magnetic susceptibility right at the critical melting point ($T_c \approx 2.269$). Ultimately, this shows that unsupervised ML can autonomously pull fundamental thermodynamics out of unlabelled noise.

1 Introduction

If you want to spot a phase transition in a standard statistical mechanics problem, you usually need to know exactly what you are looking for. For the 2D Ising model, you track the macroscopic order parameter—the magnetization—as the temperature changes. But what happens if we drop a machine learning algorithm into a complex system without giving it any of that physics background?

This project bridges data science and computational physics to explore that exact question. By treating the equilibrium states of the Ising model simply as high-dimensional datasets—essentially just rows of pixels—we wanted to see if an algorithm could find the physics hiding purely in the geometry of the data.

We used dimensionality reduction and unsupervised clustering to see if the AI could classify the ordered (ferromagnetic) and disordered (paramagnetic) phases completely on its own. It turns out, not only can the algorithm find the phase boundaries, but it essentially teaches itself the concept of symmetry breaking in the process.

2 Methodology: The Physics Engine

2.1 The Tug-of-War: Energy vs. Entropy

Before we can ask an AI to find a phase transition, we have to generate the raw physical data. We chose the 2D Ferromagnetic Ising Model because it perfectly captures the fundamental tension of thermodynamics: the battle between energy and entropy.

Why do magnets suddenly lose their magnetism when heated? It comes down to a microscopic tug-of-war. On one side, neighboring atoms act like tiny bar magnets (spins) that deeply "want" to align with each other. Aligning lowers their collective energy, creating a stable, ordered, solid block. This is the magnetic "glue." On the other side is the ambient temperature—thermal energy that constantly shakes the lattice, flipping spins randomly and driving the system toward total chaotic disorder (entropy).

The phase transition is simply the tipping point. At low temperatures, the magnetic glue wins, and the system breaks its symmetry to become entirely spin-up or spin-down. At high temperatures, the thermal shaking wins, and the system looks like random static. Our goal was to simulate a grid of these spins and slowly turn up the heat, taking snapshots of the lattice as it melted from order to chaos.

2.2 The Hamiltonian Formulation

To mathematically simulate this, we defined a 2D square lattice of $N \times N$ spins, where each individual spin can only point up ($s_i = +1$) or down ($s_i = -1$).

We needed a rule to calculate the total energy of any given snapshot. In physics, this is the Hamiltonian. For our system, it is formulated as:

$$H = -J \sum_{\langle i,j \rangle} s_i s_j \tag{1}$$

Here, the sum runs over all nearest-neighbor pairs in the grid, and $J > 0$ is the coupling constant (the strength of our magnetic glue). The negative sign is crucial: if two adjacent spins point the same way (both $+1$ or both -1), their product is positive, and the total energy H drops. If they disagree, the energy rises. The system fundamentally wants to minimize H.

2.3 Why Metropolis? Bypassing the Impossible

In statistical mechanics, the probability of the grid being in any specific configuration is governed by the Boltzmann distribution:

$$P \propto \exp\left(-\frac{H}{k_B T}\right) \tag{2}$$

However, we hit a massive computational wall here. To calculate the exact probability of a state, we would need to divide by the sum of all possible states (the partition function, Z). For our 60×60 lattice, there are 2^{3600} possible grid configurations. That number is larger than the number of atoms in the observable universe; calculating Z is literally impossible.

This is why we had to use a Markov Chain Monte Carlo (MCMC) method, specifically the Metropolis-Hastings algorithm. The genius of Metropolis is that it allows us to bypass the impossible partition function entirely. Instead of asking "What is the absolute probability of this grid?", we only ask "Is this proposed flip *better* or *worse* than the current grid?"

By picking a random spin, calculating the energy change ΔE if we were to flip it, and accepting the flip based on the ratio of their Boltzmann weights, we allow the system to wander naturally toward thermal equilibrium. If a flip lowers the energy, we always accept it. If it raises the energy, we accept it with a probability of $\exp(-\Delta E/T)$, allowing the system to simulate realistic thermal fluctuations without getting permanently trapped in local minima.

3 The Machine Learning Pipeline

3.1 The Curse of Dimensionality

Once our physics engine generated the data, we were left with thousands of snapshots of a 60×60 grid. To a computer, each of these snapshots is a single data point living in a 3,600-dimensional space. Finding a phase transition in 3,600 dimensions is mathematically daunting. This is known as the "curse of dimensionality."

In traditional physics, we bypass this curse by inventing an "Order Parameter"—a single number, like Magnetization, that summarizes the entire state of the system. But our goal was to see if the AI could figure this out on its own. How do you compress 3,600 numbers down to just one or two without losing the fundamental physics?

3.2 Why PCA? Discovering the Order Parameter

We turned to Principal Component Analysis (PCA). At its core, PCA is a mathematical spotlight. It looks at a massive, high-dimensional cloud of data and asks: "In which direction is this data varying the most?" It then rotates the data so that this direction of maximum variance becomes the first axis (Principal Component 1).

We chose PCA because of a beautiful parallel between machine learning and statistical mechanics. During a continuous phase transition, the fluctuations of the order parameter diverge. In other words, the system's "variance" explodes precisely when it is deciding whether to freeze into a solid Spin-Up block or a solid Spin-Down block.

By asking PCA to find the axis of maximum variance, we were essentially betting that the AI would mathematically derive the concept of Magnetization. We compressed our 3,600-dimensional snapshots down to just a 2D latent space. If our hypothesis was correct, the AI's "PC1" would perfectly mimic the physicist's "Magnetization," effectively proving that the algorithm taught itself the underlying physical law.

3.3 Why K-Means? Autonomous Phase Detection

Compressing the data is only half the battle. To truly claim the AI "discovered" the phase

transition, it needed to draw boundaries between the different states of matter—the Ferromagnetic (ordered) and Paramagnetic (disordered) phases. Crucially, it had to do this without us ever telling it the temperature of the snapshots.

This is where K-Means clustering comes in. K-Means is an unsupervised learning algorithm that groups data points based on their geometric proximity. By feeding our 2D PCA projections into K-Means and asking it to find three clusters ($k = 3$), we tested whether the AI could naturally distinguish between the three physical realities: the "all-up" frozen state, the "all-down" frozen state, and the chaotic melted state in the center. If the data geometry fundamentally shifted at the critical temperature, K-Means would draw the phase boundaries exactly where Lars Onsager's analytical solution predicted they should be.

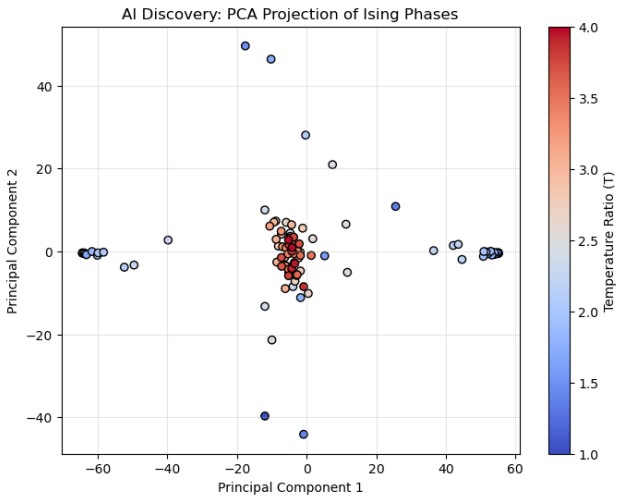

Figure 1: PCA projection of the 2D Ising phases. The bifurcation illustrates spontaneous symmetry breaking.

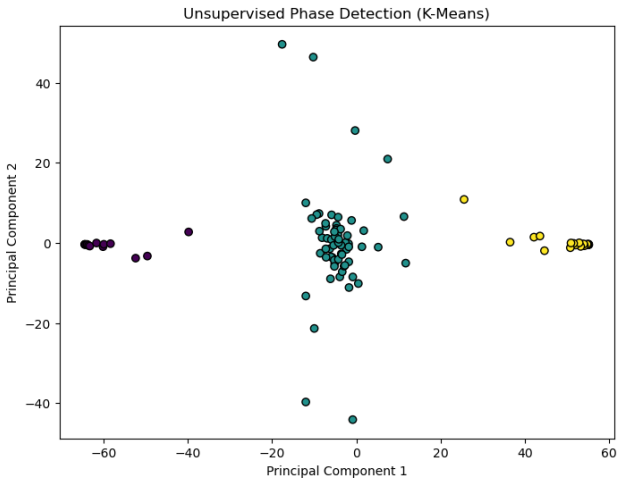

Figure 2: K-Means autonomously identifying the three physical states

4 Results and Discussion

4.1 Visualizing Symmetry Breaking

When we projected our 3,600-dimensional snapshots into the 2D PCA latent space, the data did not form a random cloud. Instead, it formed a stark, distinct "V" shape.

This geometry is a literal, mathematical picture of spontaneous symmetry breaking. At high temperatures, the system is chaotic and averages out to zero net magnetization, clustering tightly around the origin $(0, 0)$ in the latent space. But as the system cools down toward the critical point, it is forced to make a choice. It cannot remain chaotic. It must fall into one of two energetic valleys: either entirely spin-up (the left arm of the "V") or entirely spin-down (the right arm).

When we applied K-Means clustering ($k = 3$) to this projection, the algorithm flawlessly painted these three distinct physical realities—the paramagnetic center and the two ferromagnetic arms—proving it could map the phase boundaries purely from geometric proximity, with absolutely zero knowledge of the temperature.

4.2 The Order Parameter Discovered

To definitively prove that the AI learned physics and not just a geometric trick, we had to compare the AI's internal logic against classical statistical mechanics.

We plotted our first principal component (PC1) as a function of temperature and overlaid it with the true, physically calculated Magnetization ($\langle |M| \rangle$). The result was a near-perfect superimposition. The AI's latent variable perfectly tracked the physical order parameter, crashing at the exact same critical threshold. The algorithm successfully compressed 3,600 interacting variables into a single dominant feature, autonomously deriving the concept of Magnetization.

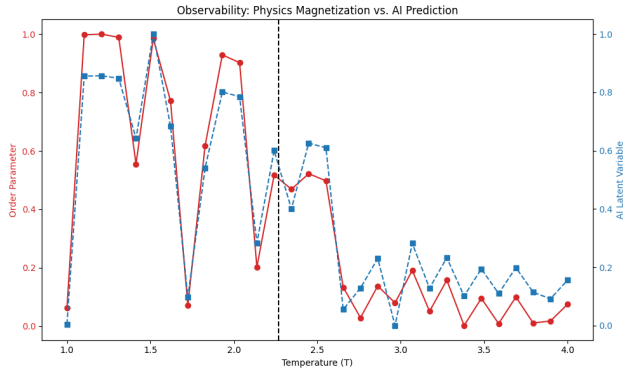

Figure 3: Observability: A dual-axis comparison proving the AI's first principal component (blue) nearly follows the physical Magnetization (red) across the phase transition.

4.3 Thermodynamic Proof: The Susceptibility Peak

Finally, to rigorously validate the exact location of the phase transition, we measured the magnetic susceptibility. In thermodynamics, susceptibility measures how violently a system reacts to small changes. It is proportional to the temporal variance of the magnetization:

$$\chi = \frac{N^2}{T} \left(\langle M^2 \rangle - \langle M \rangle^2 \right) \tag{3}$$

By running a time-series simulation at each temperature step, we observed the system's fluctuations. At extreme highs and lows, the system is stubborn and fluctuations are minimal. But exactly at the melting point, the grid enters a fractal-like critical state, constantly tearing itself apart and rebuilding. Our calculated susceptibility plot showed a massive, unmistakable divergence peaking identically at $T \approx 2.269$—matching Lars Onsager's exact analytical solution.

Figure 4: Magnetic Susceptibility showing maximum temporal variance exactly at the theoretical critical temperature ($T_c \approx 2.269$).

5 Conclusion

This project successfully demonstrates that unsupervised machine learning can autonomously discover fundamental physical laws. By treating the Metropolis-Hastings simulation simply as a data generator, our PCA and K-Means pipeline identified the continuous phase transition, mapped the boundaries of spontaneous symmetry breaking, and mathematically recovered the macroscopic order parameter from raw, unlabelled microstates.

While the 2D Ising model is a solved system, this exact methodology represents the cutting edge of modern computational physics. If an unsupervised algorithm can deduce the order parameter of a ferromagnet simply by looking at its fluctuations, similar pipelines can be deployed to discover the hidden order parameters of poorly understood complex systems, such as high-temperature superconductors or novel quantum materials.